

Memory Management Simulator

Overview

The Memory Management Simulator is a modular system designed to model the life cycle of memory in a modern operating system. It simulates physical memory allocation, multi-level cache hierarchies, and virtual memory management using paging and translation buffers aside.

Memory Layout and Assumptions

- **Physical Address Space:** 1024 bytes, organized into 16 frames.
- **Virtual Address Space:** 4096 bytes, organized into 64 pages.
- **Page/Frame Size:** Each page and frame is 64 bytes in size.
- **Alignment:** 8-byte boundary
 - Enforced in the linear allocator to simulate realistic hardware constraints
 - Prevents unaligned memory access penalties
- **Block Representation:**
 - Memory is managed as a list of blocks
 - Each block contains a header with:
 - Unique identifier (ID)
 - Allocated size
 - Requested size
 - Allocation status (free/used)
- **TLB:** 16-entry, 4-way set-associative (4 sets) caching VPN \rightarrow PFN translations using LRU
- **Page Replacement:** LRU algorithm (configurable)
- **Cache Hierarchy:** L1 (64Bytes), L2 (256Bytes), L3 (512Bytes) with LRU replacement (configurable)

Allocation Strategy Implementations

Linear Allocator (Memory Allocator)

The **Linear Allocator** manages memory using a **doubly linked list of blocks**, where each block represents a contiguous memory region that is either **allocated** or **free**.

It supports multiple allocation strategies:

- **First-Fit:** Allocates the first free block large enough for the request.
- **Best-Fit:** Allocates the smallest suitable free block to minimize internal fragmentation.
- **Worst-Fit:** Allocates the largest free block to avoid creating very small fragments.

Upon deallocation, the allocator performs **coalescing** by merging the freed block with adjacent free blocks in the list. This reduces **external fragmentation** and improves allocation efficiency.

Buddy System Design

The **Buddy System** allocator improves allocation efficiency by organizing memory into blocks whose sizes are always **powers of two** (2^k).

Power-of-Two Alignment: All memory requests are **rounded up to the next power of two**. This ensures that every allocation fits exactly into a valid buddy block size, simplifying block management and merging.

Splitting Algorithm:

If a request is made for a block of size 2^k and only a larger block is available, the allocator repeatedly **splits the larger block into two equal-sized “buddy” blocks**. This splitting continues recursively until the smallest possible power-of-two block that can satisfy the request is obtained.

Exclusive-OR (XOR) Merge Logic:

When a block is freed, the allocator computes the address of its corresponding **buddy block** using a **bitwise Exclusive-OR (XOR)** operation between the block's starting address and its size.

If the buddy block is also free, both blocks are **merged back into a single larger**

block. This merging process may continue recursively, allowing memory to be recombined into larger blocks when possible.

Cache hierarchy and replacement policy

The simulator models a **configurable three-level cache hierarchy** that closely reflects modern CPU cache design.

Each cache level varies in size, block granularity, and associativity, while the **replacement policy is selectable at runtime**, allowing flexible experimentation with different eviction strategies.

Feature	Level 1 Cache	Level 2 Cache	Level 3 Cache
Associativity	Direct-Mapped	2-Way Set-Associative	4-Way Set-Associative
Block Size	8 Bytes	16 Bytes	32 Bytes
Total Size	64 Bytes	256 Bytes	512 Bytes

Replacement policies:

```
enum ReplacementPolicy { LRU, FIFO, LFU };
```

Replacement Policy Implementations

Least Recently Used (LRU)

Evicts the cache line that has not been accessed for the longest duration.

Each cache line maintains a `last_access_time`, updated on every hit or insertion.

Eviction selects the line with the **oldest timestamp** within the set.

First-In First-Out (FIFO)

Evicts the cache line that was inserted earliest into the cache.

Each cache line records a `loaded_time` at insertion, which remains unchanged on hits.

The line with the **earliest insertion time** is selected for eviction.

Least Frequently Used (LFU):

The Least Frequently Used (LFU) replacement policy evicts the cache line that has been accessed the fewest number of times.

Each cache line maintains an access_count, which is incremented on every cache hit. When a replacement is required, the cache line with the lowest access_count is selected for eviction. If multiple cache lines have the same access frequency, a tiebreaker is applied. In such cases, the least recently used cache line—determined using its last access time—is evicted.

```
class CacheLevel {
private:
    int level_id;
    u64 size;
    u64 block_size;
    int associativity;
    u64 num_sets;
    u64 offset_bits;
    u64 index_bits;
    ReplacementPolicy policy;
    std::vector<std::vector<CacheLine>> sets;
    u64 hits = 0, misses = 0, access_counter = 0;
```

```
struct CacheLine {
    bool valid = false;
    bool dirty = false;
    u64 tag = 0;
    u64 last_access_time = 0;
    u64 insertion_time = 0;
    u64 freq = 0;
};
```

Write-Back Policy (with Write-Back Flow)

On a write operation, the cache updates only the current cache line and marks it as dirty:

```
if (is_write) line.dirty = true;|
```

No immediate write is performed to the next cache level or main memory.

The modified data is written back **only when a dirty line is evicted**:

```
if (ev_dirty) handle_writeback(ev_addr, 2);
```

Write-Back Flow:

- L1 eviction → write back to L2
- L2 eviction → write back to L3
- L3 eviction → write back to main memory (simulated)

This delayed write mechanism minimizes memory traffic and improves overall cache performance.

```
void MemoryHierarchy::handle_writeback(u64 addr, int level) {  
    if (level == 1) l2->access(addr, true);  
    else if (level == 2) l3->access(addr, true);  
}
```

Virtual memory model

The Virtual Memory subsystem translates **64-bit virtual addresses** into **physical memory addresses**. It uses a Translation Lookaside Buffer (TLB), a page table, and a page replacement mechanism to manage memory efficiently and reduce disk accesses.

Translation Lookaside Buffer (TLB):

The TLB is a **4-way set-associative cache** that stores recent virtual-to-physical page translations, reducing the overhead of page table lookups. It uses a **set associative Least Recently Used (LRU)** policy to evict the least recently accessed entry within a set.

```

struct TLBEntry {
    bool valid = false;
    u64 vpn = 0, pfn = 0, last_access = 0;
};

class TLB {
public:
    int sets, ways; u64 timer = 0;
    std::vector<std::vector<TLBEntry>> table;
    TLB(int num_entries, int assoc);
    int lookup(u64 vpn);
    void insert(u64 vpn, u64 pfn);
};

```

Page Table:

The page table is implemented as a **flat vector indexed by the Virtual Page Number (VPN)**. Each entry maintains a **Valid bit** to indicate residency in physical memory and a **Dirty bit** to track modifications. Additional metadata supports page replacement decisions.

```

struct PageTableEntry {
    bool valid = false, dirty = false, referenced = false;
    int frame_number = -1;
    u64 last_access_time = 0, loaded_time = 0;
};

```

Page Replacement Algorithms

The Virtual Memory subsystem supports multiple page replacement algorithms, selectable using the following enumeration:

```
enum PageReplacementAlgo { VM_FIFO, VM_LRU, VM_CLOCK };
```

When no page replacement algorithm is set at runtime, **LRU** is selected by default

- **VM_FIFO (First-In First-Out):**
Evicts the page that has been resident in physical memory for the longest time, based on its load order.
- **VM_LRU (Least Recently Used):**
Evicts the page with the oldest access timestamp, representing the page least recently accessed.
- **VM_CLOCK:**
Implements the Clock (Second Chance) algorithm, using a circular pointer and reference bits to approximate LRU with lower overhead.

Address Translation Flow

1. Virtual Page Number and Offset

The virtual address is divided into a Virtual Page Number (VPN) and an Offset using bitwise shifts and masks.

$$\text{Offset} = VA \& (2^k - 1)$$

$$\text{VPN} = VA \gg k$$

2. Translation Buffer Lookup

The system first searches the Translation Lookaside Buffer (TLB) for a cached VPN-to-Physical Frame Number (PFN) mapping.

3. Page Table Walk

If the TLB lookup misses, the system consults the Page Table to locate the corresponding physical frame.

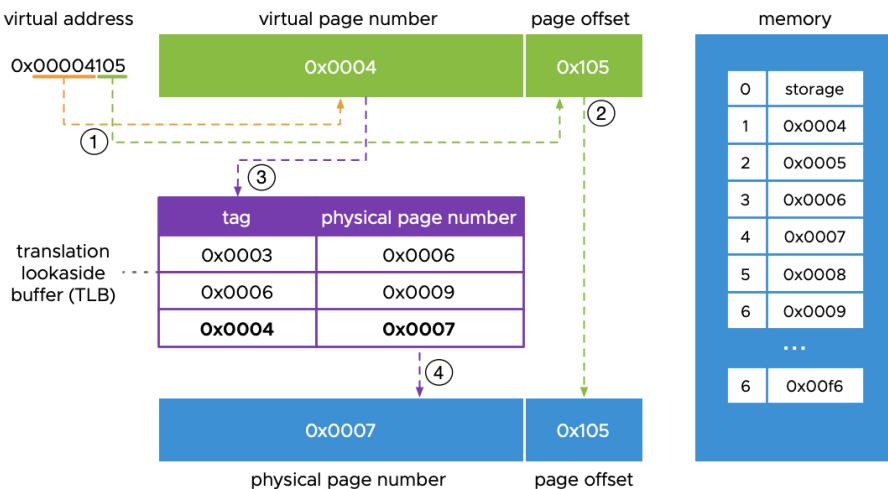
4. Physical Address Calculation

The Physical Address is computed by multiplying the Physical Frame Number by the page size and adding the offset.

$$PA = (\text{PFN} \times \text{Page Size}) + \text{Offset}$$

5. Cache System Integration

The resulting physical address is sent to the Level 1 (L1) cache to begin data access.



Limitations and simplifications

- **Fixed Cache and Virtual Memory Configuration**

Cache levels (L1, L2, L3), cache sizes, replacement policies, and virtual/physical memory sizes are fixed and small compared to real systems.

- **Simplified Memory Management Algorithms**

Allocation and page replacement use basic algorithms (Buddy, FIFO, LRU, CLOCK) without advanced techniques such as compaction, adaptive replacement, or access permissions.

- **Memory latency is ignored:** The simulator does not model cycle-accurate timing for cache or memory accesses, whereas real CPUs have precise and varying latencies for cache hits and misses across the memory hierarchy.

- **Paging Structure Limitation**

The system uses **only basic single-level (normal) paging** for virtual memory management. It does not implement more advanced paging schemes such as **hierarchical (multi-level) page tables, inverted page tables, or hashed paging**.

- **Internal Fragmentation Not Handled**

The system does not explicitly handle internal fragmentation. Memory is allocated in fixed-size pages and cache blocks, and any unused space within an allocated page or block is wasted. The simulator does not track or optimize this unused space.